

When monitoring an application running on Amazon EC2 instances, a developer might need to track system-level metrics such as memory utilization. By default, EC2 instances do not report memory usage to Amazon CloudWatch. To capture this information, the recommended approach is to install and configure the Amazon CloudWatch agent. This agent gathers additional performance data—including memory, disk, and other system metrics—and sends them to CloudWatch as custom metrics. This solution enables you to monitor the application's health and troubleshoot issues related to memory consumption.

When a developer publishes a custom metric from their Amazon EC2 instances to CloudWatch, understanding the data retention policy is crucial. CloudWatch is designed to store metric data for a specific period, ensuring that historical data remains available for analysis and troubleshooting. For custom metrics, CloudWatch retains the data for up to 15 months. This retention period allows developers to perform long-term trend analysis and evaluate the performance of their applications over time.

High-resolution metrics are important when near real-time monitoring is needed. In Amazon CloudWatch, developers can publish custom metrics with a resolution as fine as 1 second. This capability allows you to capture rapid changes in performance, which is especially useful for applications that require immediate response to performance anomalies. By using 1-second granularity, you can closely monitor the behavior of your system and react quickly to any issues that may arise.

For more granular monitoring, a developer may want to publish custom metrics with dimensions using the AWS CLI. Dimensions are key-value pairs that add context to your metrics, making it easier to filter and analyze data. The proper syntax for the `aws cloudwatch put-metric-data` command involves specifying dimensions as `key=value` pairs separated by commas. For example, using the command with dimensions like `InstanceId=i-12345abc, InstanceType=m1.small` properly structures the data for CloudWatch to process. This method enables accurate and efficient metric collection tailored to your application's needs.

Data aggregation in CloudWatch allows a developer to derive meaningful insights from raw data. When the goal is to determine the total number of requests made to an application over a specific period, the “Sum” statistic is used. The Sum aggregates the values of all individual data points, giving a clear picture of total activity. This aggregation is essential for understanding the overall load on your application and for making informed decisions about scaling and resource allocation.

Setting up cross-account monitoring in CloudWatch requires proper permissions to view metrics across different AWS accounts. The core concept here is the configuration of AWS Identity and Access Management (IAM) roles. For cross-account monitoring to work, the sharing account must have the `CloudWatch-CrossAccountSharingRole` and the monitoring account must have the `AWSServiceRoleForCloudWatchCrossAccount`. These roles allow secure access to CloudWatch data across accounts, ensuring that you can consolidate metrics for centralized monitoring without encountering access denied errors.

CloudWatch anomaly detection uses machine learning to establish what is “normal” for a given metric based on historical data. When anomaly detection is enabled, an expected range—represented as an anomaly detection band—is overlaid on your metric graphs. This band illustrates the range within which the metric values should normally fall. If the actual values deviate from this band, it signals potential issues. Understanding this expected range helps developers quickly identify anomalies and take corrective action before problems escalate.

When using CloudWatch anomaly detection, alarms are configured not on fixed thresholds but based on dynamic, data-driven boundaries. The machine learning models continuously analyze metric data and determine the expected range. If a metric value falls outside this range, an alarm is triggered. This method provides a more adaptive approach to monitoring, allowing the system to account for natural fluctuations and trends in application performance while alerting you only when truly unusual behavior is detected.

Creating an alarm in CloudWatch Metrics Insights to monitor CPU utilization involves setting a clear threshold that defines when the alarm should trigger. In this scenario, a developer wants the alarm to activate when CPU usage exceeds 80% for more than 5 consecutive minutes. The essential step is specifying the alarm threshold value, which forms the basis of the alarm's condition. By establishing this threshold, CloudWatch continuously evaluates the metric against this criterion and alerts you when the defined limit is breached, ensuring that you can promptly address potential performance issues.

When managing AWS resources via the CLI, creating alarms for metrics—such as a CPU utilization alarm—is accomplished using the `put-metric-alarm` command. This command allows you to define critical parameters, including the threshold value. In this case, by setting the threshold to 50 percent, the alarm is configured to enter the ALARM state when any instance exceeds this CPU utilization. This automated monitoring helps maintain optimal performance and triggers notifications when performance anomalies occur.

Containerized applications running on Amazon ECS require detailed monitoring for performance, resource utilization, and troubleshooting. Amazon CloudWatch Container Insights is designed specifically for this purpose. It automatically collects and aggregates metrics and logs from ECS (as well as EKS and Kubernetes environments), providing a centralized dashboard to monitor container performance. This service allows developers to analyze resource usage at a granular level, which is essential for optimizing application performance and ensuring operational excellence.

Amazon CloudWatch Lambda Insights offers valuable performance metrics for Lambda functions, including memory utilization. When a function shows high memory usage, the first step is to determine the current memory allocation. This is achieved by using the `get-function-configuration` command, which retrieves the existing configuration settings. By understanding the current allocation, a developer can make informed decisions about adjusting resources, thus optimizing function performance without overprovisioning or incurring unnecessary costs.

When an alarm indicates that CPU usage on Amazon EKS instances is exceeding thresholds, the next logical step is to investigate the underlying performance data. CloudWatch Container Insights provides detailed dashboards that display CPU usage patterns over time. By reviewing these insights, a developer can determine whether the high CPU usage is a recurring issue or merely a temporary spike. This understanding is critical before making any adjustments, such as scaling resources or modifying container configurations.

If an AWS Lambda function begins to take longer than usual to run, the first steps in troubleshooting should involve both high-level performance monitoring and detailed log analysis. Utilizing CloudWatch Lambda Insights, a developer can review performance dashboards that summarize key metrics such as function duration and memory usage. In addition, running queries in CloudWatch Logs Insights allows for a deep dive into the logs,

revealing potential issues such as errors or performance bottlenecks. These steps help diagnose the root cause of latency, ensuring that subsequent corrective actions are well informed.

For Amazon ECS environments, obtaining task-level metrics like CPU and memory utilization is crucial for effective monitoring and troubleshooting. The most efficient way to achieve this is by activating Amazon CloudWatch Container Insights on the ECS cluster. When enabled, the ECS agent automatically collects detailed performance data at the task level and sends it to CloudWatch. This built-in solution eliminates the need for custom scripts or manual instrumentation, providing a seamless and scalable monitoring experience.

When analyzing network flow data between AWS-hosted applications and on-premises destinations, it's essential to have a tool that focuses on network performance metrics. Amazon CloudWatch Network Monitor is designed specifically for this purpose. It tracks metrics such as round-trip time and packet loss for hybrid connections, giving you detailed insights into the network paths used by your data. This helps identify performance issues and network degradations, ensuring that connectivity between cloud and on-premises environments remains robust.

In scenarios where users report slow load times and intermittent outages—for example, in Europe for a global social media site—a developer needs to determine whether these issues are regional or more widespread. Amazon CloudWatch Internet Monitor provides a global view of internet performance by leveraging AWS's extensive network footprint. It collects and visualizes data about traffic patterns and connectivity health across different regions. By using Internet Monitor, the team can quickly identify if the degradation is isolated to Europe or part of a broader issue, guiding targeted performance improvements.

The Network Health Indicator (NHI) metric in Amazon CloudWatch Network Monitor is a critical measure of the network path's performance managed by AWS. When the NHI metric returns a value of 1, it indicates that there has been observed degradation in the AWS-controlled network path. This binary metric (with 1 indicating degradation and 0 indicating normal conditions) helps developers quickly ascertain whether issues in network performance exist, allowing them to focus on troubleshooting network connectivity problems.

While using Amazon CloudWatch Network Monitor, a developer might observe that the AWS Network Health Indicator for a probe shows a "Degraded" status. This status reveals two important points: first, that there is an issue with the AWS network on the monitored path, and second, that the status itself does not provide information on the operational state of the individual probe. This means the degradation is attributed to the network conditions rather than the probe malfunctioning. Such insights help developers pinpoint network issues and understand that further investigation into the probe is not necessary unless additional symptoms are observed.

On the Probe details page of Amazon CloudWatch Network Monitor, developers have the ability to manage probe configurations to ensure optimal network performance monitoring. Specifically, they can modify the probe settings—such as activating or deactivating the probe—and manage tags by adding or removing them. These configuration options allow for easy resource management and organization. However, it is important to note that performance metrics like round-trip time and packet loss are automatically measured by the system and cannot be manually adjusted. This separation between configuration and measurement ensures that the monitoring data remains accurate and reliable.